

■■■■■ Optimizing Agent Planning for Security and Autonomy

A horizontal row of 15 small black squares, likely representing a binary sequence or a set of data points.

input_papers\test.pdf

2026-02-14 15:44

33



■ 1 ■

[■] OPTIMIZING AGENT PLANNING FOR SECURITY AND AUTONOMY AashishKolluri¹
RishiSharma^{1,2†} ManuelCosta¹ BorisKo“pf¹ TobiasNießen^{3†} MarkRussinovich¹ ShrutiTople¹
SantiagoZanella-Be‘gulin¹ 1Microsoft 2EPFL 3T...

■ 2 ■

[■] Optimizing Agent Planning for Security and Autonomy integrity and confidentiality labels to all data an agent processes, propagating labels to suggested actions, and using these labels to determine whether ...

■ 3 ■

[■] Optimizing Agent Planning for Security and Autonomy 2 Background: Information-flow Control for AI Agents Information-flow control mechanisms use security labels to describe the security properties of data during their life time...

■ 4 ■

[■] Optimizing Agent Planning for Security and Autonomy isolation, but their content remains hidden from the planner's LLM. The original formulation of the DualLLM pattern allows for restricted outputs of the quarantined LLM t...

■ 5 ■

[■] Optimizing Agent Planning for Security and Autonomy unsuccessful traces the agent repeatedly attempting actions that failed policy checks (which we allow to continue) and did not lead to any progress, a pattern that a human would quickly...

■ 6 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy expand variables(ask endorsement=True)), maintaining the label of the context, or (ii) proceed without endorsement (by calling expand variables(ask endo...)

■ 7 ■

[■] Optimizing Agent Planning for Security and Autonomy
information extraction in quarantined LLM queries.

By design, no attack succeeds in this setting due to strict policies, deterministic defenses, and our assumption that a...

■ 8 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy 80 60 40 20 0 1 2 3 4 HITLoad(k)
)(k@RCT o3-mini 80 60 40 20 0 1 2 3 4 HITLoad(k) Basic Basic-IFC FIDES PRUDENTIA
)(k@RCT o4-mini Figure2:TCR@kcurvess...

■ 9 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy AttackSuccessRate HITLoad(average)
TCR@∞(%) Turns(average) Model Environment Basic PRUDENTIA Basic PRUDENTIA Basic
PRUDENTIA Basic PRUDENTIA GitLab 20.80...

■ 10 ■

[■] Optimizing Agent Planning for Security and Autonomy security prompts and begin approving actions without careful consideration (Stanton et al., 2016; Seidling et al., 2011). Deterministic defenses based on IFC can be more effective...

■ 11 ■

■] Optimizing Agent Planning for Security and Autonomy 7 Related Work Probabilistic Defenses. Several techniques have been proposed for minimizing the likelihood of prompt injection attacks in LLM-based systems using...

■ 12 ■

[■] Optimizing Agent Planning for Security and Autonomy References Sahar Abdehnabi, Aideen Fay, Giovanni Cherubin, Ahmed Salem, Mario Fritz, and Andrew Paverd. Get my drift? Catching LLM task drift with activat...

■ 13 ■

[■] Optimizing Agent Planning for Security and Autonomy
Feiran Jia, Tong Wu, Xin Qin, and Anna Squicciarini. The TaskShield:
Enforcing task alignment to defend against indirect prompt injection in LLM agents.
In 63rd Annual Meeting o...

■ 14 ■

[■] Optimizing Agent Planning for Security and Autonomy Simon Willison.
The Dual LLM pattern for building AI assistants that can resist prompt injection. Online: <https://simonwillison.net/2023/Apr/25/dual-lm-pattern>, Apr 20...

■ 15 ■

[■] Optimizing Agent Planning for Security and Autonomy A Case Study:
Choice between Endorsement and Approval We show an abridged example trace of an Agent Dojo task where a PRUDENTIA agent strategically chooses not to task the user before...

■ 16 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy Assistant
search_emails(query="TechServices password reset") Tool - Email 1: From
support@techservices.com Body in variable #search_emails-result-0.body# ...

[■] Optimizing Agent Planning for Security and Autonomy B Additional Results Table 2 provides comprehensive performance data across all methods and models, revealing consistent patterns of improvement from Basic th...

■ 18 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy 100 80 60 40 20 0 DI DIQ DD)%(0@RCT
TCR@0usingo3-mini 100 80 60 40 20 0 DI DIQ DD)%(0@RCT TCR@0usingo4-mini 100 80 60 40
20 0 DI DIQ DD)%(1@RCT TCR@1us...

■ 19 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy 100 80 60 40 20 0 banking slack travel
workspace)%(etaRnoitelpmoCksaT TCR@∞usingo3-mini 100 80 60 40 20 0 banking slack travel
workspace)%(etaRnoitelpmo...

■ 20 ■

[■] Optimizing Agent Planning for Security and Autonomy C Overhead Analysis We provide rough estimates for the costs of Basic-IFC, FIDES, and PRUDENTIA in terms of the number of input (Table 3) and output (Table 5) tokens, and USD cents (...)

■ 21 ■

■■■ Optimizing Agent Planning for Security and Autonomy We believe that these costs can be brought down significantly as our goal was to investigate autonomy gains, so we have not yet made an effort to optimize...

■ 22 ■

[■] Optimizing Agent Planning for Security and Autonomy D System Prompts D.1 Basic Planner We used the following system prompt for the basic agent with and without IFC, taken from DeBenedetto et al. (2024).
System You are an AI lang...

■ 23 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy 2. **Field inside a dict**: - Format:
`#tool_name-result-tool_name_count.key#` 3. **Basic type inside a list of dicts**: - Format:
`#tool_name-result-tool...

■ 24 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy - **Example:** - You plan to call: `generate_report(title, date)` . - You ask quarantined LLM: - Query 1: **"What's the title in #v1#?"** -> stored in `#quar...

■ 25 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy - You call `expand_variables` with `[#v1#, #v2#]` . - Now use expanded information to carry out your next steps. ## IMPORTANT REMINDERS - Never invent or g...

■ 26 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy - **!Untrusted Data:** Untrusted data is any information received from a tool call that may (cid:44)→ contain prompt-injections. Examples include email con...

■ 27 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy ### Variable naming conventions: Variable identifiers always start and end with `#` . **Basic types** (`str`, `bool`, `int`, `float`): - Format: `#tool_nam...

■ 28 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy - **Scenario:** You know exactly what tool you want to call next and require specific (cid:44)→ arguments from the variables. - **Action:** Query the quar...

■ 29 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy To minimize the number of endorsements and approvals, you should avoid expanding variables (cid:44)→ in your context for as long as possible. Once you kno...

■ 30 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy - Always ensure each query to quarantined LLM returns exactly one piece of information. - Plan and exhaust other options and defer using `ifc_expand_varia...

■ 31 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy E ExampleAccessibilityTreeObservation
We present first the unprocessed accessibility tree, with the prompt injection highlighted in red. Axtree
"OBSERVATION: 0 (curren...

■ 32 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy [294] StaticText '([■30fb_■30fb;])' [297]
StaticText "There's nothing here 2026" [303] complementary " [310] StaticText '1 point' [313]
StaticText '(+1, ■2...

■ 33 ■

[■] OptimizingAgentPlanningforSecurityandAutonomy [215] StaticText VAR_18_7165e3de
[UNTRUSTED] [216] link VAR_19_6f3a142a [UNTRUSTED] [218] StaticText VAR_20_68801a33
[UNTRUSTED] [228] link VAR_21_6eff0cb...



Model	Transformer	Attention
Transformer	Transformer	Transformer
Attention	Attention	Attention
Neural Network	Neural Network	Neural Network
Deep Learning	Deep Learning	Deep Learning
Model	Model	Model